

# sampling-and-cv

October 16, 2023

```
[1]: # Loading Libraries
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
import numpy as np
```

```
[2]: # Dataframe
df = pd.read_csv('Data.csv')
df.head()
```

```
[2]: Unnamed: 0  gender  SeniorCitizen  Partner  Dependents  tenure  \
0          6607        1              0         0           1         1
1          2598        0              0         0           0         7
2          2345        0              0         0           1         4
3          4093        0              0         0           0        29
4           693        0              0         0           0         3

    PhoneService  MultipleLines  InternetService  OnlineSecurity  ...  \
0              0              1              0              0  ...
1              1              0              1              0  ...
2              1              0              2              1  ...
3              1              2              1              0  ...
4              1              2              1              0  ...

    DeviceProtection  TechSupport  StreamingTV  StreamingMovies  Contract  \
0                  0            0            0              0          0
1                  2            0            0              0          0
2                  1            1            1              1          0
3                  0            0            0              0          0
4                  0            0            0              0          0

    PaperlessBilling  PaymentMethod  MonthlyCharges  TotalCharges  Churn
0                  1              2          25.30         2153      1
1                  1              2          75.15         4396      0
2                  1              0          20.05         6211      0
3                  1              1          76.00         1850      0
4                  1              1          75.10         2350      1
```

[5 rows x 21 columns]

```
[3]: # Drop unnecessary columns/features
df.drop('Unnamed: 0',axis=1, inplace=True)
df.head()
```

```
[3]:  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0      1              0        0             1         1             0
1      0              0        0             0         7             1
2      0              0        0             1         4             1
3      0              0        0             0        29             1
4      0              0        0             0         3             1

      MultipleLines  InternetService  OnlineSecurity  OnlineBackup  \
0                1                0                0                0
1                0                1                0                0
2                0                2                1                1
3                2                1                0                0
4                2                1                0                0

      DeviceProtection  TechSupport  StreamingTV  StreamingMovies  Contract  \
0                    0            0            0                0          0
1                    2            0            0                0          0
2                    1            1            1                1          0
3                    0            0            0                0          0
4                    0            0            0                0          0

      PaperlessBilling  PaymentMethod  MonthlyCharges  TotalCharges  Churn
0                    1              2          25.30         2153      1
1                    1              2          75.15         4396      0
2                    1              0          20.05         6211      0
3                    1              1          76.00         1850      0
4                    1              1          75.10         2350      1
```

```
[4]: df.shape
```

```
[4]: (5282, 20)
```

```
[5]: # Count of Unique Values
df.Churn.value_counts()
```

```
[5]: 0    3892
1    1390
Name: Churn, dtype: int64
```

```
[6]: # No value percentage
1390 / (3892+1390)
```

```
[6]: 0.2631578947368421
```

```
[7]: # Summary of Dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5282 entries, 0 to 5281
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 5282 non-null   int64
1   SeniorCitizen          5282 non-null   int64
2   Partner                5282 non-null   int64
3   Dependents             5282 non-null   int64
4   tenure                 5282 non-null   int64
5   PhoneService           5282 non-null   int64
6   MultipleLines          5282 non-null   int64
7   InternetService        5282 non-null   int64
8   OnlineSecurity         5282 non-null   int64
9   OnlineBackup           5282 non-null   int64
10  DeviceProtection       5282 non-null   int64
11  TechSupport            5282 non-null   int64
12  StreamingTV            5282 non-null   int64
13  StreamingMovies        5282 non-null   int64
14  Contract               5282 non-null   int64
15  PaperlessBilling       5282 non-null   int64
16  PaymentMethod          5282 non-null   int64
17  MonthlyCharges         5282 non-null   float64
18  TotalCharges           5282 non-null   int64
19  Churn                  5282 non-null   int64
dtypes: float64(1), int64(19)
memory usage: 825.4 KB
```

## 0.1 Analysis the dataset using seaborn / matplotlib / scatter

```
[8]: # Correlation between features
df.corr()
```

```
[8]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure \
gender	1.000000	-0.005691	-0.010143	0.011058	0.001303
SeniorCitizen	-0.005691	1.000000	0.016648	-0.211271	0.006176
Partner	-0.010143	0.016648	1.000000	0.447629	0.382432
Dependents	0.011058	-0.211271	0.447629	1.000000	0.162933
tenure	0.001303	0.006176	0.382432	0.162933	1.000000

PhoneService	-0.009706	0.000581	0.031578	-0.005830	0.005912
MultipleLines	-0.000430	0.129040	0.151068	-0.020722	0.354790
InternetService	-0.005294	-0.040186	0.007205	0.041972	-0.032037
OnlineSecurity	-0.020189	-0.129719	0.159928	0.157262	0.325451
OnlineBackup	-0.025773	-0.012166	0.166939	0.095640	0.371342
DeviceProtection	-0.009305	-0.015922	0.168815	0.076783	0.369331
TechSupport	-0.010619	-0.151078	0.132582	0.137274	0.326993
StreamingTV	-0.012994	0.024338	0.136715	0.045948	0.288135
StreamingMovies	-0.011420	0.043124	0.122310	0.023907	0.301600
Contract	0.000555	-0.151939	0.303243	0.243080	0.671184
PaperlessBilling	-0.014090	0.156417	-0.020634	-0.109935	0.004043
PaymentMethod	0.010188	-0.035050	-0.160535	-0.040414	-0.360323
MonthlyCharges	-0.018822	0.219945	0.105603	-0.114920	0.253605
TotalCharges	-0.022718	0.040956	0.069859	-0.013196	0.152843
Churn	-0.011997	0.146549	-0.150053	-0.164490	-0.345544

	PhoneService	MultipleLines	InternetService	\
gender	-0.009706	-0.000430	-0.005294	
SeniorCitizen	0.000581	0.129040	-0.040186	
Partner	0.031578	0.151068	0.007205	
Dependents	-0.005830	-0.020722	0.041972	
tenure	0.005912	0.354790	-0.032037	
PhoneService	1.000000	-0.016345	0.385682	
MultipleLines	-0.016345	1.000000	-0.105796	
InternetService	0.385682	-0.105796	1.000000	
OnlineSecurity	-0.007874	0.006028	-0.027201	
OnlineBackup	0.017500	0.134460	0.032058	
DeviceProtection	0.000422	0.132798	0.048320	
TechSupport	-0.006139	0.019266	-0.018510	
StreamingTV	0.055390	0.172380	0.101060	
StreamingMovies	0.048362	0.188860	0.092672	
Contract	0.005342	0.120023	0.097158	
PaperlessBilling	0.014489	0.172369	-0.141856	
PaymentMethod	0.006362	-0.176313	0.095054	
MonthlyCharges	0.247419	0.436398	-0.325588	
TotalCharges	0.081045	0.118773	-0.064748	
Churn	0.010122	0.031270	-0.048820	

	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	\
gender	-0.020189	-0.025773	-0.009305	-0.010619	
SeniorCitizen	-0.129719	-0.012166	-0.015922	-0.151078	
Partner	0.159928	0.166939	0.168815	0.132582	
Dependents	0.157262	0.095640	0.076783	0.137274	
tenure	0.325451	0.371342	0.369331	0.326993	
PhoneService	-0.007874	0.017500	0.000422	-0.006139	
MultipleLines	0.006028	0.134460	0.132798	0.019266	
InternetService	-0.027201	0.032058	0.048320	-0.018510	

OnlineSecurity	1.000000	0.186626	0.182355	0.276510
OnlineBackup	0.186626	1.000000	0.191348	0.189892
DeviceProtection	0.182355	0.191348	1.000000	0.247866
TechSupport	0.276510	0.189892	0.247866	1.000000
StreamingTV	0.057760	0.142882	0.278896	0.174178
StreamingMovies	0.065996	0.151145	0.301894	0.172420
Contract	0.371159	0.282646	0.352138	0.428750
PaperlessBilling	-0.146473	-0.019611	-0.032091	-0.113245
PaymentMethod	-0.085892	-0.111508	-0.135513	-0.097672
MonthlyCharges	-0.049605	0.125340	0.162808	-0.001710
TotalCharges	0.036493	0.090223	0.103120	0.053532
Churn	-0.288926	-0.193152	-0.173138	-0.274718

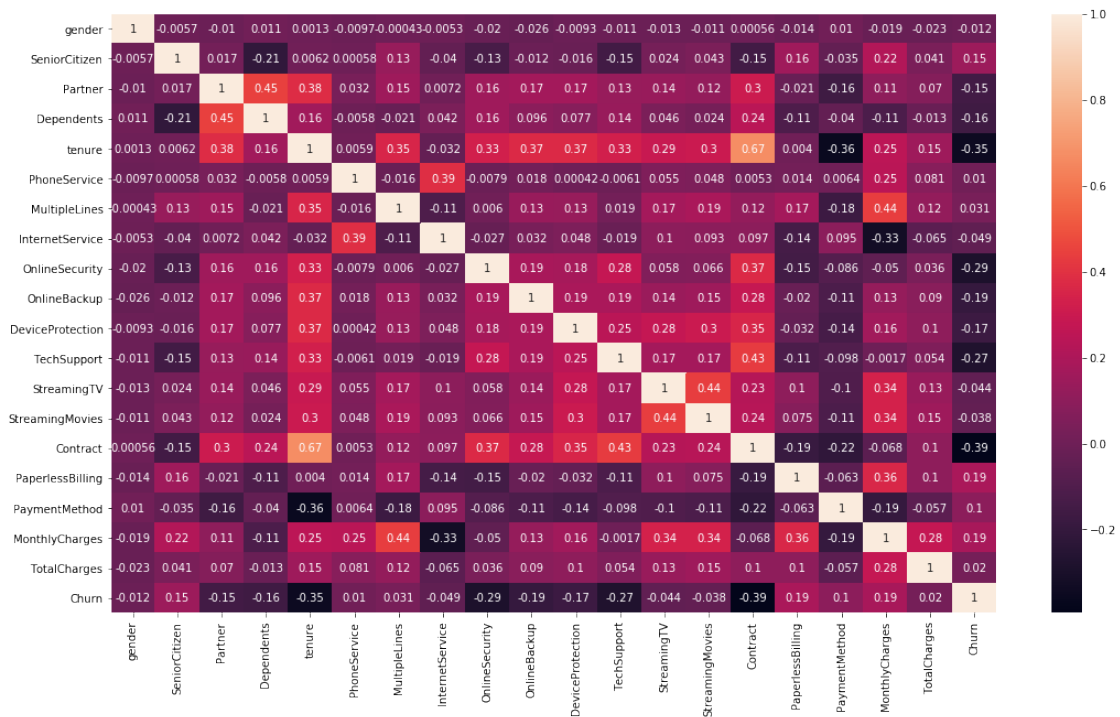
	StreamingTV	StreamingMovies	Contract	PaperlessBilling	\
gender	-0.012994	-0.011420	0.000555	-0.014090	
SeniorCitizen	0.024338	0.043124	-0.151939	0.156417	
Partner	0.136715	0.122310	0.303243	-0.020634	
Dependents	0.045948	0.023907	0.243080	-0.109935	
tenure	0.288135	0.301600	0.671184	0.004043	
PhoneService	0.055390	0.048362	0.005342	0.014489	
MultipleLines	0.172380	0.188860	0.120023	0.172369	
InternetService	0.101060	0.092672	0.097158	-0.141856	
OnlineSecurity	0.057760	0.065996	0.371159	-0.146473	
OnlineBackup	0.142882	0.151145	0.282646	-0.019611	
DeviceProtection	0.278896	0.301894	0.352138	-0.032091	
TechSupport	0.174178	0.172420	0.428750	-0.113245	
StreamingTV	1.000000	0.437809	0.231143	0.101389	
StreamingMovies	0.437809	1.000000	0.236128	0.075255	
Contract	0.231143	0.236128	1.000000	-0.185507	
PaperlessBilling	0.101389	0.075255	-0.185507	1.000000	
PaymentMethod	-0.100597	-0.114956	-0.218531	-0.063408	
MonthlyCharges	0.338557	0.339162	-0.067540	0.359566	
TotalCharges	0.134112	0.150553	0.104879	0.101619	
Churn	-0.043920	-0.038240	-0.394490	0.188793	

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
gender	0.010188	-0.018822	-0.022718	-0.011997
SeniorCitizen	-0.035050	0.219945	0.040956	0.146549
Partner	-0.160535	0.105603	0.069859	-0.150053
Dependents	-0.040414	-0.114920	-0.013196	-0.164490
tenure	-0.360323	0.253605	0.152843	-0.345544
PhoneService	0.006362	0.247419	0.081045	0.010122
MultipleLines	-0.176313	0.436398	0.118773	0.031270
InternetService	0.095054	-0.325588	-0.064748	-0.048820
OnlineSecurity	-0.085892	-0.049605	0.036493	-0.288926
OnlineBackup	-0.111508	0.125340	0.090223	-0.193152
DeviceProtection	-0.135513	0.162808	0.103120	-0.173138

TechSupport	-0.097672	-0.001710	0.053532	-0.274718
StreamingTV	-0.100597	0.338557	0.134112	-0.043920
StreamingMovies	-0.114956	0.339162	0.150553	-0.038240
Contract	-0.218531	-0.067540	0.104879	-0.394490
PaperlessBilling	-0.063408	0.359566	0.101619	0.188793
PaymentMethod	1.000000	-0.194857	-0.056972	0.100015
MonthlyCharges	-0.194857	1.000000	0.279822	0.186615
TotalCharges	-0.056972	0.279822	1.000000	0.020294
Churn	0.100015	0.186615	0.020294	1.000000


```
[9]: plt.figure(figsize=(18,10))
sns.heatmap(df.corr(),annot=True)
```

```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x26417db08c8>
```




```
[10]: sns.pairplot(df)
```

```
[10]: <seaborn.axisgrid.PairGrid at 0x264186384c8>
```




```
[11]: sns.countplot(df.gender,hue='Churn',data=df)
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x2642b64fb48>
```



```
[12]: sns.countplot(df['SeniorCitizen'], hue=df.Churn)
```


```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x2642b509d08>
```






```
[13]: sns.countplot(df.Dependents,hue=df.Churn)
```

```
[13]: <matplotlib.axes._subplots.AxesSubplot at 0x2642bdc2748>
```




```
[14]: sns.countplot(df.PhoneService,hue=df.Churn)
```

```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0x2642f3d3348>
```




```
[15]: sns.countplot(df['InternetService'], hue=df.Churn)
```

```
[15]: <matplotlib.axes._subplots.AxesSubplot at 0x2642f3d3588>
```




```
[16]: sns.countplot(df.TechSupport,hue="Churn" , data=df)
```

```
[16]: <matplotlib.axes._subplots.AxesSubplot at 0x26430429508>
```




```
[17]: sns.barplot(df.Churn,df.TechSupport,hue='gender',data=df)
```

```
[17]: <matplotlib.axes._subplots.AxesSubplot at 0x2643045dc08>
```



```
[18]: sns.catplot(x="Churn", y="MonthlyCharges", hue="gender", data=df)
```

```
[18]: <seaborn.axisgrid.FacetGrid at 0x2643042cc48>
```



```
[19]: x = df.drop('Churn',axis=1)
      y = df.Churn
```

```
[20]: x.shape
```

```
[20]: (5282, 19)
```

```
[21]: y.value_counts()
```

```
[21]: 0    3892
      1    1390
      Name: Churn, dtype: int64
```

# 1 ML model and evaluating model by cross validation before sampling

## 1.1 Hold out Cross Validation

```
[22]: from sklearn.model_selection import train_test_split
      xtrain, xtest, ytrain, ytest=train_test_split(x,y, test_size=0.3,
      ↪random_state=42)
```

```
[23]: xtrain.shape
```

```
[23]: (3697, 19)
```

```
[24]: ytrain.shape
```

```
[24]: (3697,)
```

### 1.1.1 Decision Tree Classifier

```
[25]: from sklearn.tree import DecisionTreeClassifier
      dtc = DecisionTreeClassifier()
```

```
[26]: dtc.fit(xtrain,ytrain)
```

```
[26]: DecisionTreeClassifier()
```

```
[27]: score_dtc = dtc.score(xtest,ytest)
```

```
[28]: score_dtc
```

```
[28]: 0.7129337539432177
```

### 1.1.2 Random Forest Classifier

```
[29]: from sklearn.ensemble import RandomForestClassifier
      rfc = RandomForestClassifier()
```

```
[30]: rfc.fit(xtrain,ytrain)
```

```
[30]: RandomForestClassifier()
```

```
[31]: score_rfc = rfc.score(xtest,ytest)
```

```
[32]: score_rfc
```

```
[32]: 0.7899053627760252
```

### 1.1.3 Xstream Gradient boosting (XGBoost) Classifier

```
[33]: pip install xgboost
```

```
Requirement already satisfied: xgboost in d:\anaconda3\lib\site-packages (1.6.1)  
Requirement already satisfied: numpy in d:\anaconda3\lib\site-packages (from  
xgboost) (1.18.1)  
Requirement already satisfied: scipy in d:\anaconda3\lib\site-packages (from  
xgboost) (1.4.1)  
Note: you may need to restart the kernel to use updated packages.
```

```
[34]: from xgboost import XGBClassifier  
xgbc = XGBClassifier()
```

```
[35]: xgbc.fit(xtrain,ytrain)
```

```
[35]: XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,  
                    colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,  
                    early_stopping_rounds=None, enable_categorical=False,  
                    eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',  
                    importance_type=None, interaction_constraints='',  
                    learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,  
                    max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,  
                    missing=nan, monotone_constraints='()', n_estimators=100,  
                    n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=0,  
                    reg_alpha=0, reg_lambda=1, ...)
```

```
[36]: score_xgbc = xgbc.score(xtest,ytest)
```

```
[37]: score_xgbc
```

```
[37]: 0.7728706624605678
```

### 1.1.4 AdaBoost Classifier

```
[38]: from sklearn.ensemble import AdaBoostClassifier  
abc = AdaBoostClassifier()
```

```
[39]: abc.fit(xtrain,ytrain)
```

```
[39]: AdaBoostClassifier()
```

```
[40]: score_abc = abc.score(xtest,ytest)
```

```
[41]: score_abc
```

```
[41]: 0.7867507886435331
```

### 1.1.5 K-Nearest Neighbour (KNN) Classifier

```
[42]: from sklearn.neighbors import KNeighborsClassifier  
      knnc = KNeighborsClassifier()
```

```
[43]: knnc.fit(xtrain,ytrain)
```

```
[43]: KNeighborsClassifier()
```

```
[44]: score_knnc = knnc.score(xtest,ytest)
```

```
[45]: score_knnc
```

```
[45]: 0.7488958990536277
```

### 1.1.6 Logistic Regression

```
[46]: from sklearn.linear_model import LogisticRegression  
      lrc = LogisticRegression()
```

```
[47]: lrc.fit(xtrain,ytrain)
```

D:\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
[47]: LogisticRegression()
```

```
[48]: score_lrc = lrc.score(xtest,ytest)
```

```
[49]: score_lrc
```

```
[49]: 0.7886435331230284
```

## 1.2 Comparison

```
[50]: cdf = pd.DataFrame(['Decision Tree',score_dtc],['Random_  
↳Forest',score_rfc],['XGBoost',score_xgbc],['Ada_  
↳Boost',score_abc],['K-Nearest Neighbor',score_knnc],['Logistic_  
↳Regression',score_lrc]],columns=['Classifier', 'Accuracy'])
```



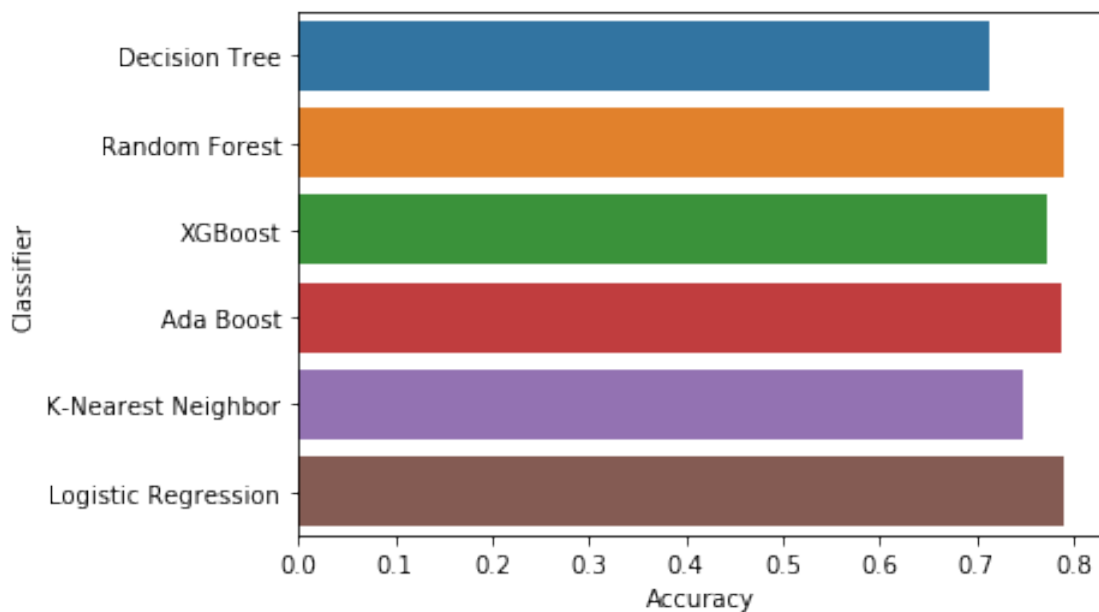
```
cdf
```

```
[50]:
```

	Classifier	Accuracy
0	Decision Tree	0.712934
1	Random Forest	0.789905
2	XGBoost	0.772871
3	Ada Boost	0.786751
4	K-Nearest Neighbor	0.748896
5	Logistic Regression	0.788644

```
[51]: sns.barplot(cdf.Accuracy, cdf.Classifier)
```

```
[51]: <matplotlib.axes._subplots.AxesSubplot at 0x26431f2b048>
```



### 1.2.1 K-Fold Cross Validation

```
[52]: from sklearn.model_selection import KFold, cross_val_score
kf = KFold(n_splits=5) #each fold will contain 20% data
```

```
[53]: result_kf = cross_val_score(lrc, x, y, cv=kf)
```

D:\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,  
D:\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,  
D:\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,  
D:\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,  
D:\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

[54]: result\_kf

```
[54]: array([0.81740776, 0.78807947, 0.7907197 , 0.77556818, 0.80681818])
```

```
[55]: result_kf.min()
```

```
[55]: 0.7755681818181818
```

```
[56]: result_kf.max()
```

```
[56]: 0.8174077578051088
```

### 1.2.2 Stratified k-fold Cross Validation

```
[57]: from sklearn.model_selection import StratifiedKFold  
skf = StratifiedKFold()
```

```
[58]: result_skf = cross_val_score(lrc, x, y, cv=skf)
```

```
D:\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
D:\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,  
D:\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
```

```
D:\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:818:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
```

```
D:\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:818:
```

```
ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
```

```
[59]: result_skf
```

```
[59]: array([0.81078524, 0.79564806, 0.79166667, 0.77840909, 0.79734848])
```

```
[60]: result_skf.max()
```

```
[60]: 0.8107852412488175
```

```
[61]: result_skf.min()
```

```
[61]: 0.7784090909090909
```

### 1.2.3 Leave One Out Cross Validation(LOOC)

```
[62]: #from sklearn.model_selection import LeaveOneOut
      #loo = LeaveOneOut()
      #result = cross_val_score(lrc,X,y,cv=loo)
      #result
      #result.mean()
```

## 2 sampling on given dataset and Create ML model and evaluating model by cross validation again (after sampling)

```
[63]: y.value_counts()
```

```
[63]: 0    3892  
      1    1390  
      Name: Churn, dtype: int64
```

```
[64]: x.shape
```

```
[64]: (5282, 19)
```

### 2.1 Sampling

### 2.2 S1. CROSS Validation with Synthetic Minority Oversampling Technique (SMOTETomek)

```
[65]: pip install imblearn
```

```
Collecting imblearn  
  Using cached imblearn-0.0-py2.py3-none-any.whl (1.9 kB)  
Collecting imbalanced-learn  
  Using cached imbalanced_learn-0.9.1-py3-none-any.whl (199 kB)  
Requirement already satisfied: threadpoolctl>=2.0.0 in d:\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (3.1.0)  
  Using cached imbalanced_learn-0.9.0-py3-none-any.whl (199 kB)  
Requirement already satisfied: scikit-learn>=1.0.1 in d:\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.0.2)  
Requirement already satisfied: scipy>=1.1.0 in d:\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.4.1)  
Requirement already satisfied: numpy>=1.14.6 in d:\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.18.1)  
Requirement already satisfied: joblib>=0.11 in d:\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (0.14.1)  
Installing collected packages: imbalanced-learn, imblearn  
Successfully installed imbalanced-learn-0.9.0 imblearn-0.0  
Note: you may need to restart the kernel to use updated packages.
```

```
[66]: from imblearn.combine import SMOTETomek  
      smt = SMOTETomek()
```

```
[70]: x_smt, y_smt = smt.fit_resample(x,y)
```

```
[71]: x_smt.shape
```

```
[71]: (7190, 19)
```

```
[73]: y_smt.value_counts()
```

```
[73]: 1    3595  
      0    3595  
      Name: Churn, dtype: int64
```

## 2.3 Hold out cross validation

```
[74]: x_smt_train,x_smt_test,y_smt_train,y_smt_test =  
      ↪train_test_split(x_smt,y_smt,train_size=.7,random_state=42)
```

```
[76]: x_smt_train.shape
```

```
[76]: (5033, 19)
```

```
[77]: x_smt_test.shape
```

```
[77]: (2157, 19)
```

### 2.3.1 Decision Tree Classifier

```
[78]: dtc_smt = DecisionTreeClassifier()  
      dtc_smt.fit(x_smt_train,y_smt_train)
```

```
[78]: DecisionTreeClassifier()
```

```
[79]: score_dtc_smt = dtc_smt.score(x_smt_test,y_smt_test)
```

```
[80]: score_dtc_smt
```

```
[80]: 0.8191933240611962
```

### 2.3.2 Random Forest Classifier

```
[81]: rfc_smt = RandomForestClassifier()  
      rfc_smt.fit(x_smt_train,y_smt_train)
```

```
[81]: RandomForestClassifier()
```

```
[82]: score_rfc_smt = rfc_smt.score(x_smt_test,y_smt_test)  
      score_rfc_smt
```

```
[82]: 0.8687992582290218
```

### 2.3.3 XGBoost Classifier

```
[83]: xgbc_smt = XGBClassifier()  
xgbc_smt.fit(x_smt_train,y_smt_train)
```

```
[83]: XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,  
                  colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,  
                  early_stopping_rounds=None, enable_categorical=False,  
                  eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',  
                  importance_type=None, interaction_constraints='',  
                  learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,  
                  max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,  
                  missing=nan, monotone_constraints='()', n_estimators=100,  
                  n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=0,  
                  reg_alpha=0, reg_lambda=1, ...)
```

```
[84]: score_xgbc_smt = xgbc_smt.score(x_smt_test,y_smt_test)  
score_xgbc_smt
```

```
[84]: 0.866481223922114
```

### 2.3.4 Ada Boost classifier

```
[85]: abc_smt = AdaBoostClassifier()  
abc_smt.fit(x_smt_train,y_smt_train)
```

```
[85]: AdaBoostClassifier()
```

```
[86]: score_abc_smt = abc_smt.score(x_smt_test,y_smt_test)  
score_abc_smt
```

```
[86]: 0.8437644877144181
```

### 2.3.5 KNN Classifier

```
[87]: knnc_smt = KNeighborsClassifier()  
knnc_smt.fit(x_smt_train,y_smt_train)
```

```
[87]: KNeighborsClassifier()
```

```
[88]: score_knnc_smt = knnc_smt.score(x_smt_test,y_smt_test)  
score_knnc_smt
```

```
[88]: 0.799721835883171
```

### 2.3.6 Logistic Regression

```
[89]: lrc_smt = LogisticRegression()  
lrc_smt.fit(x_smt_train,y_smt_train)
```

```
D:\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
[89]: LogisticRegression()
```

```
[90]: score_lrc_smt = lrc_smt.score(x_smt_test,y_smt_test)  
score_lrc_smt
```

```
[90]: 0.8293926750115902
```

### 2.3.7 Comparison

```
[182]: df_smt = pd.DataFrame([['Decision Tree (SMOTETomek)',score_dtc_smt],['Random_↵  
↵Forest (SMOTETomek)',score_rfc_smt],['XGBoost_↵  
↵(SMOTETomek)',score_xgbc_smt],['Ada Boost_↵  
↵(SMOTETomek)',score_abc_smt],['K-Nearest Neighbor_↵  
↵(SMOTETomek)',score_knnc_smt],['Logistic Regression_↵  
↵(SMOTETomek)',score_lrc_smt]],columns=['Classifier', 'Accuracy'])  
df_smt
```

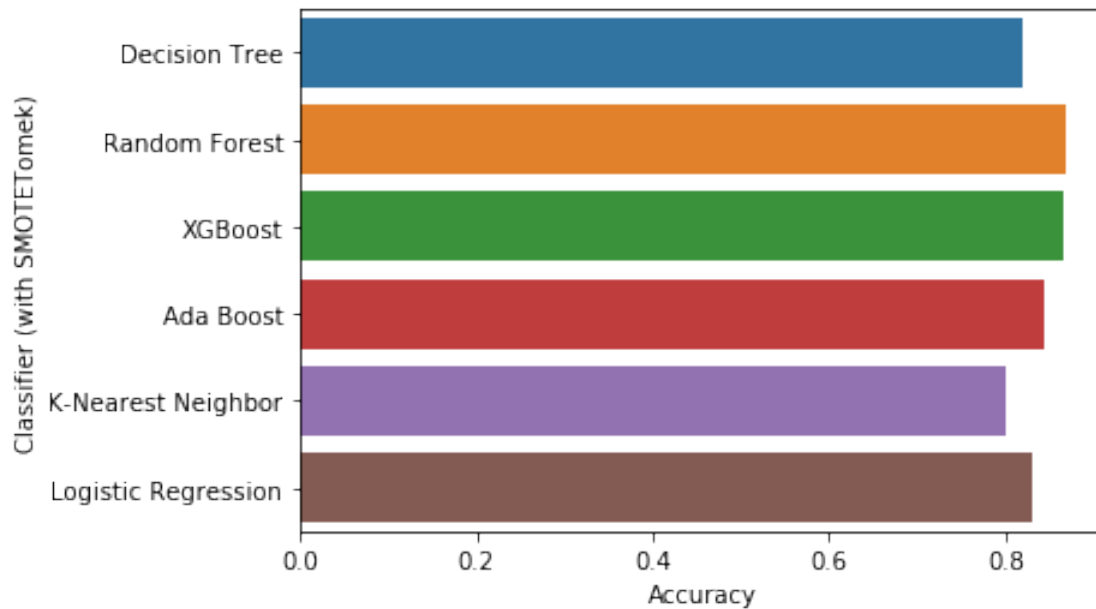
```
[182]:
```

	Classifier	Accuracy
0	Decision Tree (SMOTETomek)	0.819193
1	Random Forest (SMOTETomek)	0.868799
2	XGBoost (SMOTETomek)	0.866481
3	Ada Boost (SMOTETomek)	0.843764
4	K-Nearest Neighbor (SMOTETomek)	0.799722
5	Logistic Regression (SMOTETomek)	0.829393

```
[94]: sns.barplot(df_smt.Accuracy,df_smt.Classifier])
```

```
[94]: <matplotlib.axes._subplots.AxesSubplot at 0x26430558a48>
```





### 2.3.8 K-Fold Cross Validation

```
[95]: kf_smt = KFold(n_splits=5)
      result_kf_smt = cross_val_score(rfc_smt, x_smt, y_smt, cv=kf_smt)
      result_kf_smt
```

```
[95]: array([0.80667594, 0.81641168, 0.79763561, 0.91655076, 0.95201669])
```

```
[96]: result_kf.min()
```

```
[96]: 0.7755681818181818
```

```
[97]: result_kf.max()
```

```
[97]: 0.8174077578051088
```

### 2.3.9 Stratified k-fold Cross Validation

```
[98]: skf_smt = StratifiedKFold()
      result_skf_smt = cross_val_score(rfc_smt, x_smt, y_smt, cv=skf_smt)
      result_skf_smt
```

```
[98]: array([0.73504868, 0.80945758, 0.91655076, 0.89916551, 0.90403338])
```

```
[99]: result_skf_smt.max()
```

```
[99]: 0.9165507649513213
```

```
[100]: result_skf_smt.min()
```

```
[100]: 0.7350486787204451
```

### 2.3.10 Leave-One-Out-Cross-Validation(LOOC)

```
[102]: #loo_smt = LeaveOneOut()  
#result_smt = cross_val_score(rfc_smt,x_smt,y_smt,cv=loo_smt)  
#result_smt  
#result_smt.mean()
```

## 3 S2. Cross Validation with Near Miss

```
[104]: from imblearn.under_sampling import NearMiss  
nm = NearMiss()
```

```
[106]: x_nm,y_nm = nm.fit_resample(x,y)
```

```
[109]: y.value_counts()
```

```
[109]: 0    3892  
      1    1390  
      Name: Churn, dtype: int64
```

```
[108]: y_nm.value_counts()
```

```
[108]: 1    1390  
      0    1390  
      Name: Churn, dtype: int64
```

### 3.1 Hold out Cross Validation

```
[110]: x_nm_train,x_nm_test,y_nm_train,y_nm_test =  
      ↪train_test_split(x_nm,y_nm,train_size=.7,random_state=42)
```

```
[111]: x_nm_train.shape
```

```
[111]: (1945, 19)
```

```
[112]: x_nm_test.shape
```

```
[112]: (835, 19)
```

### 3.1.1 Decision Tree Classifier

```
[113]: dtc_nm = DecisionTreeClassifier()  
dtc_nm.fit(x_nm_train,y_nm_train)
```

```
[113]: DecisionTreeClassifier()
```

```
[114]: score_dtc_nm = dtc_nm.score(x_nm_test,y_nm_test)
```

```
[115]: score_dtc_nm
```

```
[115]: 0.629940119760479
```

### 3.1.2 Random Forest Classifier

```
[116]: rfc_nm = RandomForestClassifier()  
rfc_nm.fit(x_nm_train,y_nm_train)
```

```
[116]: RandomForestClassifier()
```

```
[117]: score_rfc_nm = rfc_nm.score(x_nm_test,y_nm_test)  
score_rfc_nm
```

```
[117]: 0.6574850299401198
```

### 3.1.3 XGBoost Classifier

```
[118]: xgbc_nm = XGBClassifier()  
xgbc_nm.fit(x_nm_train,y_nm_train)
```

```
[118]: XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,  
                  colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,  
                  early_stopping_rounds=None, enable_categorical=False,  
                  eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',  
                  importance_type=None, interaction_constraints='',  
                  learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,  
                  max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,  
                  missing=nan, monotone_constraints='()', n_estimators=100,  
                  n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=0,  
                  reg_alpha=0, reg_lambda=1, ...)
```

```
[119]: score_xgbc_nm = xgbc_nm.score(x_nm_test,y_nm_test)  
score_xgbc_nm
```

```
[119]: 0.6730538922155689
```

### 3.1.4 Ada Boost classifier

```
[120]: abc_nm = AdaBoostClassifier()  
abc_nm.fit(x_nm_train,y_nm_train)
```

```
[120]: AdaBoostClassifier()
```

```
[121]: score_abc_nm = abc_nm.score(x_nm_test,y_nm_test)  
score_abc_nm
```

```
[121]: 0.6778443113772455
```

### 3.1.5 KNN Classifier

```
[122]: knnc_nm = KNeighborsClassifier()  
knnc_nm.fit(x_nm_train,y_nm_train)
```

```
[122]: KNeighborsClassifier()
```

```
[123]: score_knnc_nm = knnc_nm.score(x_nm_test,y_nm_test)  
score_knnc_nm
```

```
[123]: 0.578443113772455
```

### 3.1.6 Logistic Regression

```
[124]: lrc_nm = LogisticRegression()  
lrc_nm.fit(x_nm_train,y_nm_train)
```

```
[124]: LogisticRegression()
```

```
[125]: score_lrc_nm = lrc_nm.score(x_nm_test,y_nm_test)  
score_lrc_nm
```

```
[125]: 0.6634730538922156
```

### 3.1.7 Comparison

```
[185]: df_nm = pd.DataFrame(['Decision Tree (Near Miss)',score_dtc_nm],['Random_█  
↳Forest (Near Miss)',score_rfc_nm],['XGBoost (Near_█  
↳Miss)',score_xgbc_nm],['Ada Boost (Near Miss)',score_abc_nm],['K-Nearest_█  
↳Neighbor (Near Miss)',score_knnc_nm],['Logistic Regression (Near_█  
↳Miss)',score_lrc_nm]],columns=['Classifier', 'Accuracy'])  
df_nm
```

```
[185]:
```

	Classifier	Accuracy
0	Decision Tree (Near Miss)	0.629940


```

1      Random Forest (Near Miss)  0.657485
2      XGBoost (Near Miss)  0.673054
3      Ada Boost (Near Miss)  0.677844
4      K-Nearest Neighbor (Near Miss)  0.578443
5      Logistic Regression (Near Miss)  0.663473

```

```
[174]: sns.barplot(df_nm.Accuracy,df_nm.Classifier)
```

```
[174]: <matplotlib.axes._subplots.AxesSubplot at 0x2643474cb08>
```



### 3.1.8 K-Fold Cross Validation

```
[130]: kf_nm = KFold(n_splits=5)
result_kf_nm = cross_val_score(abc_nm, x_nm, y_nm,cv=kf_nm)
result_kf_nm
```

```
[130]: array([0.26978417, 0.4676259 , 0.72841727, 0.52697842, 0.47661871])
```

```
[131]: result_kf_nm.min()
```

```
[131]: 0.2697841726618705
```

```
[132]: result_kf_nm.max()
```

```
[132]: 0.7284172661870504
```

### 3.1.9 Stratified k-fold Cross Validation

```
[133]: skf_nm = StratifiedKFold()  
result_skf_nm = cross_val_score(abc_nm, x_nm, y_nm, cv=skf_nm)  
result_skf_nm
```

```
[133]: array([0.50179856, 0.64568345, 0.67625899, 0.69064748, 0.69244604])
```

```
[134]: result_skf_nm.min()
```

```
[134]: 0.5017985611510791
```

```
[135]: result_skf_nm.max()
```

```
[135]: 0.6924460431654677
```

### 3.1.10 Leave-One-Out-Cross-Validation(LOOC)

```
[136]: #loo_nm = LeaveOneOut()  
#result_nm = cross_val_score(abc_nm, x_nm, y_nm, cv=loo_nm)  
#result_nm  
#result_nm.mean()
```

## 4 S3. Cross Validation with Random Over Sampler

```
[138]: from imblearn.over_sampling import RandomOverSampler  
ros = RandomOverSampler()
```

```
[139]: x_ros, y_ros = ros.fit_resample(x, y)
```

```
[140]: y.value_counts()
```

```
[140]: 0    3892  
      1    1390  
      Name: Churn, dtype: int64
```

```
[141]: y_ros.value_counts()
```

```
[141]: 1    3892  
      0    3892  
      Name: Churn, dtype: int64
```

#### 4.0.1 Hold Out Cross Validation

```
[142]: x_ros_train,x_ros_test,y_ros_train,y_ros_test =  
↳train_test_split(x_ros,y_ros,train_size=.7,random_state=42)
```

```
[145]: x_ros_train.shape
```

```
[145]: (5448, 19)
```

```
[146]:
```

```
[146]: (2336, 19)
```

#### 4.0.2 Decision Tree Classifier

```
[147]: dtc_ros = DecisionTreeClassifier()  
dtc_ros.fit(x_ros_train,y_ros_train)
```

```
[147]: DecisionTreeClassifier()
```

```
[149]: score_dtc_ros = dtc_ros.score(x_ros_test,y_ros_test)  
score_dtc_ros
```

```
[149]: 0.8608732876712328
```

#### 4.0.3 Random Forest Classifier

```
[150]: rfc_ros = RandomForestClassifier()  
rfc_ros.fit(x_ros_train,y_ros_train)
```

```
[150]: RandomForestClassifier()
```

```
[151]: score_rfc_ros = rfc_ros.score(x_ros_test,y_ros_test)  
score_rfc_ros
```

```
[151]: 0.875
```

#### 4.0.4 XGBoost Classifier

```
[152]: xgbc_ros = XGBClassifier()  
xgbc_ros.fit(x_ros_train,y_ros_train)
```

```
[152]: XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,  
colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,  
early_stopping_rounds=None, enable_categorical=False,  
eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',  
importance_type=None, interaction_constraints='',
```

```
learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,  
max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,  
missing=nan, monotone_constraints='()', n_estimators=100,  
n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=0,  
reg_alpha=0, reg_lambda=1, ...)
```

```
[153]: score_xgbc_ros = xgbc_ros.score(x_ros_test,y_ros_test)  
score_xgbc_ros
```

```
[153]: 0.8441780821917808
```

#### 4.0.5 Ada Boost classifier

```
[154]: abc_ros = AdaBoostClassifier()  
abc_ros.fit(x_ros_train,y_ros_train)
```

```
[154]: AdaBoostClassifier()
```

```
[155]: score_abc_ros = abc_ros.score(x_ros_test,y_ros_test)  
score_abc_ros
```

```
[155]: 0.7551369863013698
```

#### 4.0.6 KNN Classifier

```
[156]: knnc_ros = KNeighborsClassifier()  
knnc_ros.fit(x_ros_train,y_ros_train)
```

```
[156]: KNeighborsClassifier()
```

```
[157]: score_knnc_ros = knnc_ros.score(x_ros_test,y_ros_test)  
score_knnc_ros
```

```
[157]: 0.7277397260273972
```

#### 4.0.7 Logistic Regression

```
[158]: lrc_ros = LogisticRegression()  
lrc_ros.fit(x_ros_train,y_ros_train)
```

D:\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:



```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-  
regression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
```

```
[158]: LogisticRegression()
```

```
[159]: score_lrc_ros = lrc_ros.score(x_ros_test,y_ros_test)  
score_lrc_ros
```

```
[159]: 0.7478595890410958
```

#### 4.0.8 Comparison


```
[184]: df_ros = pd.DataFrame(['Decision Tree (Random Over_  
↳Sampler)',score_dtc_ros],['Random Forest (Random Over_  
↳Sampler)',score_rfc_ros],['XGBoost (Random Over_  
↳Sampler)',score_xgbc_ros],['Ada Boost (Random Over_  
↳Sampler)',score_abc_ros],['K-Nearest Neighbor (Random Over_  
↳Sampler)',score_knnc_ros],['Logistic Regression (Random Over_  
↳Sampler)',score_lrc_ros]],columns=['Classifier', 'Accuracy'])  
df_ros
```

```
[184]:
```

	Classifier	Accuracy
0	Decision Tree (Random Over Sampler)	0.860873
1	Random Forest (Random Over Sampler)	0.875000
2	XGBoost (Random Over Sampler)	0.844178
3	Ada Boost (Random Over Sampler)	0.755137
4	K-Nearest Neighbor (Random Over Sampler)	0.727740
5	Logistic Regression (Random Over Sampler)	0.747860

```
[162]: sns.barplot(df_ros.Accuracy,df_ros.Classifier)
```

```
[162]: <matplotlib.axes._subplots.AxesSubplot at 0x264348341c8>
```



#### 4.0.9 K-Fold Cross Validation

```
[163]: kf_ros = KFold(n_splits=5)
result_kf_ros = cross_val_score(abc_ros, x_ros, y_ros,cv=kf_ros)
result_kf_ros
```

```
[163]: array([0.72254335, 0.72575466, 0.71612075, 0.77007065, 0.68187661])
```

```
[164]: result_kf_ros.min()
```

```
[164]: 0.6818766066838047
```

```
[165]: result_kf_ros.max()
```

```
[165]: 0.7700706486833655
```

#### 4.0.10 Stratified k-fold Cross Validation

```
[166]: skf_ros = StratifiedKFold()
result_skf_ros = cross_val_score(abc_ros, x_ros, y_ros,cv=skf_ros)
result_skf_ros
```

```
[166]: array([0.77520873, 0.77135517, 0.77777778, 0.7495183 , 0.77763496])
```

```
[167]: result_skf_ros.min()
```

```
[167]: 0.7495183044315993
```

```
[168]: result_skf_ros.max()
```

```
[168]: 0.7777777777777778
```

#### 4.0.11 Leave-One-Out-Cross-Validation(LOOC)

```
[169]: #loo_ros = LeaveOneOut()  
#result_ros = cross_val_score(abc_ros,x_ros,y_ros,cv=loo_ros)  
#result_ros  
#result_ros.mean()
```

## 5 Evaluate all results

```
[200]: df_compare = pd.concat([cdf,df_smt,df_nm,df_ros])  
df_compare
```

```
[200]:
```

	Classifier	Accuracy
0	Decision Tree	0.712934
1	Random Forest	0.789905
2	XGBoost	0.772871
3	Ada Boost	0.786751
4	K-Nearest Neighbor	0.748896
5	Logistic Regression	0.788644
0	Decision Tree (SMOTETomek)	0.819193
1	Random Forest (SMOTETomek)	0.868799
2	XGBoost (SMOTETomek)	0.866481
3	Ada Boost (SMOTETomek)	0.843764
4	K-Nearest Neighbor (SMOTETomek)	0.799722
5	Logistic Regression (SMOTETomek)	0.829393
0	Decision Tree (Near Miss)	0.629940
1	Random Forest (Near Miss)	0.657485
2	XGBoost (Near Miss)	0.673054
3	Ada Boost (Near Miss)	0.677844
4	K-Nearest Neighbor (Near Miss)	0.578443
5	Logistic Regression (Near Miss)	0.663473
0	Decision Tree (Random Over Sampler)	0.860873
1	Random Forest (Random Over Sampler)	0.875000
2	XGBoost (Random Over Sampler)	0.844178
3	Ada Boost (Random Over Sampler)	0.755137
4	K-Nearest Neighbor (Random Over Sampler)	0.727740
5	Logistic Regression (Random Over Sampler)	0.747860

```
[201]: df_compare.reset_index(drop=True,inplace=True)
```

```
[202]: df_compare
```

```
[202]:
```

	Classifier	Accuracy
0	Decision Tree	0.712934
1	Random Forest	0.789905
2	XGBoost	0.772871
3	Ada Boost	0.786751
4	K-Nearest Neighbor	0.748896
5	Logistic Regression	0.788644
6	Decision Tree (SMOTETomek)	0.819193
7	Random Forest (SMOTETomek)	0.868799
8	XGBoost (SMOTETomek)	0.866481
9	Ada Boost (SMOTETomek)	0.843764
10	K-Nearest Neighbor (SMOTETomek)	0.799722
11	Logistic Regression (SMOTETomek)	0.829393
12	Decision Tree (Near Miss)	0.629940
13	Random Forest (Near Miss)	0.657485
14	XGBoost (Near Miss)	0.673054
15	Ada Boost (Near Miss)	0.677844
16	K-Nearest Neighbor (Near Miss)	0.578443
17	Logistic Regression (Near Miss)	0.663473
18	Decision Tree (Random Over Sampler)	0.860873
19	Random Forest (Random Over Sampler)	0.875000
20	XGBoost (Random Over Sampler)	0.844178
21	Ada Boost (Random Over Sampler)	0.755137
22	K-Nearest Neighbor (Random Over Sampler)	0.727740
23	Logistic Regression (Random Over Sampler)	0.747860

```
[195]: df_compare.Accuracy.max()
```

```
[195]: 0.875
```

```
[206]: # Locating Maximum row values
df_compare.loc[df_compare['Accuracy'].idxmax()]
```

```
[206]: Classifier    Random Forest (Random Over Sampler)
Accuracy                                0.875
Name: 19, dtype: object
```

```
[193]: plt.figure(figsize=(10,10))
sns.barplot(x='Accuracy',y='Classifier',data=df_compare)
```

```
[193]: <matplotlib.axes._subplots.AxesSubplot at 0x26436e3c948>
```

